

Report: Heap Data Structures - Implementation, Analysis, and Applications

Author: Raghul Krishnan

Date: 01-26-2025

Heapsort Implementation and Analysis

1. Implementation

Heapsort is implemented using a max-heap. Key steps include:

1. Building a max-heap from the input array.
2. Repeatedly extracting the maximum element and placing it at the end of the array.
3. Maintaining the heap property after each extraction.

[GitHub implementation link](#)

2. Analysis

- **Time Complexity:**
 - **Building the heap:** $O(n)$. The heapify operation runs in $O(\log n)$ time, and it is called for $n/2$ elements during heap construction.
 - **Extracting elements:** Each extraction requires $O(\log n)$, and this is performed n times.
 - Total complexity: $O(n) + O(n \log n) = O(n \log n)$.
 - **Best, Average, and Worst Case:** $O(n \log n)$ in all cases because heap operations do not depend on input order.
- **Space Complexity:**
 - $O(1)$ additional space since sorting is done in-place. No extra data structures are required beyond the input array.
- **Overhead:**
 - The overhead mainly comes from recursive or iterative calls to maintain the heap property during heapify.

3. Comparison with Quicksort and Merge Sort

- Running Heapsort, Quicksort, and Merge Sort on datasets of varying sizes and distributions (sorted, reverse-sorted, random) shows:
 - Quicksort is often faster on average for random data due to better cache locality.
 - Merge Sort provides consistent performance and excels in stable sorting scenarios.
 - Heapsort performs reliably but can be slower than Quicksort due to additional comparisons during heap operations.
 - **Theoretical Analysis:**
 - Quicksort: Average $O(n \log n)$, Worst $O(n^2)$.
 - Merge Sort: $O(n \log n)$ in all cases, $O(n)$ space complexity.
 - Heapsort: Consistent $O(n \log n)$ with $O(1)$ space.
-

Priority Queue Implementation and Applications

Part A: Priority Queue Implementation

[GitHub implementation link](#)

1. Data Structure:

- Used a Python list to represent the binary heap. This allows for efficient implementation of heap operations via array indexing.
- Task objects are used to represent individual tasks, including fields like task ID, priority, and other metadata.
- A max-heap is chosen as it suits applications like scheduling, where the highest priority task should be processed first.

2. Core Operations:

- **insert(task):**
 - Adds a new task to the heap, places it at the end, and adjusts the heap property by "bubbling up."
 - **Time Complexity:** $O(\log n)$.

- **extract_max():**
 - Removes and returns the task with the highest priority (root of the heap). The last element replaces the root, and the heap property is restored by "bubbling down."
 - **Time Complexity:** $O(\log n)$.
- **increase_key(task, new_priority):**
 - Updates a task's priority and ensures the heap property by bubbling up if the priority increases or down if it decreases.
 - **Time Complexity:** $O(\log n)$.
- **is_empty():**
 - Simple check to see if the heap is empty.
 - **Time Complexity:** $O(1)$.

3. Scheduling Simulation:

- Tasks are dynamically added and processed based on priority, demonstrating real-world scenarios like CPU scheduling or event management.
- Observed behavior aligns with theoretical expectations; higher-priority tasks are consistently processed before lower-priority ones.

Design Choices and Observations

1. **Heap Representation:** A list-based heap is simple and efficient due to the direct mapping of parent-child relationships in a binary tree to array indices.
2. **Task Priority:** Using custom objects allows flexibility in storing additional metadata while prioritizing operations efficiently.
3. **Max-Heap:** Ideal for scenarios requiring the highest-priority task first, e.g., event management or job scheduling.

Conclusion

- Heapsort provides a reliable, in-place sorting solution with consistent $O(n \log n)$ performance. However, its real-world efficiency often lags behind Quicksort due to higher overheads.

- Priority queues, implemented using a heap, are versatile tools for efficiently managing dynamic, prioritized data. They find applications in real-time systems, networking, and more.
- The empirical comparison and analysis highlight the trade-offs between sorting and priority queue operations, emphasizing the need to choose the right tool for specific scenarios.